

# **Project Progress Report**

**By: Jeremy Flagg**

**Spring Semester 2025**

## **1. Study Overview**

This study focuses on:

- Training Merjek AI models on a GPU cluster.

## **2. Early Steps & Prompt Generation**

The initial phase involved testing different LLM models for prompt generation and analyzing their outputs after database insertion.

## **3. Models Tested**

Several models were tested for effectiveness and performance:

- Open-source models (e.g., LLama 3.1 8B, DeepSeek R1 1.5B, Mistral 7B v0.3)

## **4. GPU Cluster Specifications**

Cluster Quota specifications:

- Max Jobs: 6
- Max Nodes: 3
- Max GPUs per Job: 4
- Max Runtime per Job: 48 hours

**Training Progress:(1/24)**

- Initial meeting

**Training Progress: (2/7)**

- Installation of Ollama and different open-source LLM models.
- Prompt generation and insertion into MySQL Workbench.

**Training Progress:(2/14-2/28)**

- Initial training/test practice, locally and in GPU Cluster,with Human Trafficking and Campus csv files.
- Migration to MongoDB Atlas/Compass

**Training Progress: (3/7/25)**

- Dataset: 2,000 documents (subset of 10K)
- Split: 80% train, 20% test
- Tested on 2 GPUs (1 node)
- Estimated training time: ~52 minutes for 1 epoch

### **Training Progress: (3/14/25)**

- Created Merjek Github
- Prompt generation 5 hours per 1000 documents
- Meeting at library helping Md with MongoDB setup and prompt generation
- Continue generating ~8K prompts for the entire dataset of ~10K documents.
- Mistral 7B v0.3 is the model used for prompt generation. (LM Studio on my Windows setup)
- After generation, iterated through MongoDB collection to add prompts into arrays.
- Edit Slurm training script before executing within GPU cluster.
- Scaled training from 2,000 docs at 1 epoch to 10,000 docs at 3 epochs.

 **Loaded 305835 valid prompts from the first 10,000 documents.**

**Training samples: 244668**

**Validation samples: 61167**

 **Using device: cuda, Batch size: 16**

**GPU #: 4**

**Estimated train time for 1 epoch: 4 hours 41 minutes**

**Estimated train time for 3 epochs: 14 hours 4 minutes**

## View inside cluster after 1 epoch for 10K docs:

```
PS C:\WINDOWS\system32> ssh jmflagg@itiger.memphis.edu
jmflagg@itiger.memphis.edu's password:
Last login: Tue Mar 11 20:17:26 2025 from 10.228.110.243
[jmflagg@itiger ~]$ cd /project/jmflagg/merjek-study/
[jmflagg@itiger merjek-study]$ squeue -u $USER
        JOBID PARTITION     NAME     USER ST       TIME  NODES MODELLIST(REASON)
        5421    bigTiger merjekai    jmflagg R      2:59:00      1 itiger04
[jmflagg@itiger merjek-study]$ tail merjekai-training-output.txt
{'eval_loss': 8.805888175964355, 'eval_runtime': 43.2928, 'eval_samples_per_second': 1412.866, 'eval_steps_per_second': 22.082, 'epoch': 0.62}
{'loss': 8.8116, 'grad_norm': 208776.609375, 'learning_rate': 7.496730316505363e-06, 'epoch': 0.63}
{'eval_loss': 8.804204940795898, 'eval_runtime': 42.7269, 'eval_samples_per_second': 1431.582, 'eval_steps_per_second': 22.375, 'epoch': 0.63}
{'loss': 8.8209, 'grad_norm': 193645.765625, 'learning_rate': 7.444415380591159e-06, 'epoch': 0.63}
{'eval_loss': 8.803586959838867, 'eval_runtime': 42.5937, 'eval_samples_per_second': 1436.057, 'eval_steps_per_second': 22.445, 'epoch': 0.63}
{'loss': 8.8142, 'grad_norm': 205195.5625, 'learning_rate': 7.392100444676957e-06, 'epoch': 0.63}
{'eval_loss': 8.803701400756836, 'eval_runtime': 43.2752, 'eval_samples_per_second': 1413.442, 'eval_steps_per_second': 22.091, 'epoch': 0.63}
{'loss': 8.8328, 'grad_norm': 210443.34375, 'learning_rate': 7.339785508762752e-06, 'epoch': 0.63}
{'eval_loss': 8.803175926208496, 'eval_runtime': 43.3248, 'eval_samples_per_second': 1411.825, 'eval_steps_per_second': 22.066, 'epoch': 0.63}
{'loss': 8.7729, 'grad_norm': 209517.65625, 'learning_rate': 7.287470572848549e-06, 'epoch': 0.64}
[jmflagg@itiger merjek-study]$ exit
logout
Connection to itiger.memphis.edu closed.
PS C:\WINDOWS\system32> ssh jmflagg@itiger.memphis.edu
jmflagg@itiger.memphis.edu's password:
Last login: Tue Mar 11 21:17:54 2025 from 10.228.110.238
[jmflagg@itiger ~]$ cd /project/jmflagg/merjek-study/
[jmflagg@itiger merjek-study]$ tail merjekai-training-output.txt
{'loss': 8.7498, 'grad_norm': 210646.0, 'learning_rate': 6.800941668846455e-08, 'epoch': 1.0}
{'eval_loss': 8.748345375061035, 'eval_runtime': 42.6307, 'eval_samples_per_second': 1434.811, 'eval_steps_per_second': 22.425, 'epoch': 1.0}
{'loss': 8.7193, 'grad_norm': 204133.09375, 'learning_rate': 1.5694480774261054e-08, 'epoch': 1.0}
{'eval_loss': 8.74834156036377, 'eval_runtime': 42.8912, 'eval_samples_per_second': 1426.097, 'eval_steps_per_second': 22.289, 'epoch': 1.0}
{'train_runtime': 16880.9389, 'train_samples_per_second': 14.494, 'train_steps_per_second': 0.226, 'train_loss': 8.908902582622092, 'epoch': 1.0}
Evaluating model...
Evaluation results: {'eval_loss': 8.74834156036377, 'eval_runtime': 42.6337, 'eval_samples_per_second': 1434.71, 'eval_steps_per_second': 22.424, 'epoch': 1.0}
Saving model to ./fine-tuned-model-merjekai3
[✓] Model and tokenizer saved successfully.
[✓] Training job completed.
```

## View inside cluster after 3 epochs for 10K:

```
[jmflagg@itiger ~]$ cd /project/jmflagg/merjek-study/
[jmflagg@itiger merjek-study]$ head merjekai-training-output.txt
Starting merjekai.py...
Starting merjekai.py...
[✓] Connected to MongoDB Atlas successfully.
[✓] Loaded 305835 valid prompts from the first 10,000 documents.
Training samples: 244668
Validation samples: 61167
Using device: cuda, Batch size: 16
Starting training...
{'loss': 9.2462, 'grad_norm': 176753.75, 'learning_rate': 1.99825616880286e-05, 'epoch': 0.0}
{'eval_loss': 9.237972259521484, 'eval_runtime': 43.31, 'eval_samples_per_second': 1412.307, 'eval_steps_per_second': 22.073, 'epoch': 0.0}
[jmflagg@itiger merjek-study]$ tail merjekai-training-output.txt
{'loss': 8.2178, 'grad_norm': 242138.140625, 'learning_rate': 3.3132792745662224e-08, 'epoch': 3.0}
{'eval_loss': 8.314220428466797, 'eval_runtime': 43.37, 'eval_samples_per_second': 1410.351, 'eval_steps_per_second': 22.043, 'epoch': 3.0}
{'loss': 8.2782, 'grad_norm': 253271.640625, 'learning_rate': 1.5694480774261054e-08, 'epoch': 3.0}
{'eval_loss': 8.314230918884277, 'eval_runtime': 43.1083, 'eval_samples_per_second': 1418.914, 'eval_steps_per_second': 22.177, 'epoch': 3.0}
{'train_runtime': 50719.8742, 'train_samples_per_second': 14.472, 'train_steps_per_second': 0.226, 'train_loss': 8.52074397049928, 'epoch': 3.0}
Evaluating model...
Evaluation results: {'eval_loss': 8.31423282623291, 'eval_runtime': 42.679, 'eval_samples_per_second': 1433.188, 'eval_steps_per_second': 22.4, 'epoch': 3.0}
Saving model to ./fine-tuned-model-merjekai3
[✓] Model and tokenizer saved successfully.
[✓] Training job completed.
[jmflagg@itiger merjek-study]$
```

## View inside MongoDB Compass:

Documents10.1K

Aggregations

Schema


Indexes1


Validation


{Label:10000}

Generat

+ ADD DATA ▾

 EXPORT DATA ▾

 UPDATE

 DELETE

```
_id: ObjectId('67be90d2e152ac3375cc4939')
Label : 10000
Url : "https://www.memphis.edu/gradschool/resources/graduate_faculty/cas/swrk..."
Title : "Social Work Graduate Faculty Resources -
        Graduate School
        - The Un..."
Text : "

        Social Work Graduate Faculty Resources -
        Graduate School
        ..."
Client : Object
Prompts : Array (30)
  0: "university memphis campus"
  1: " university memphis academic calendar"
  2: " university memphis admissions process"
  3: " university of memphis faculty members"
  4: " university memphis degrees"
  5: " university memphis school of social work"
  6: " university of memphis application deadline"
  7: " university of memphis faculty positions"
  8: " university memphis doctoral programs"
  9: " university of memphis external graduate faculty"
  10: " university of memphis financial aid"
```

**Training Progress: (3/21/25 - 3/28/25)**

**Goal: Find a way to utilize the GPU cluster for prompt generation**

#### **1. Initial Attempt with vLLM**

- **Tried using vLLM for running LLaMA 3.1 8B.**
  - **Faced challenges and decided to move on to other methods.**
- 

#### **2. Transformers Library Approach**

- **Attempted to use the Transformers library and download LLaMA 3.1 8B from HuggingFace.**
  - **Performance was poor, comparable to 1B models.**
  - **Assumed Ollama optimizes models behind the scenes for better performance.**
- 

#### **3. Dockerized Ollama Installation**

- **Installed a Dockerized version of Ollama on the cluster using Podman (compatible with Docker).**
  - **Downloaded and tested the LLaMA 3.2B model.**
  - **The model worked but encountered two major issues:**
    - **No External Access: Unable to access Ollama from outside the container, even though ports were open and listening.**
    - **CPU-Only Inference: Without external access, couldn't create a Slurm script for GPU usage, resulting in CPU-only inference.**
- 

#### **4. Native Ollama Installation**

- **Installed Ollama natively by downloading and extracting the binary to a directory.**

- Still faced CPU-only inference since GPUs on the cluster are only accessible through Slurm.
- 

## 5. Understanding Cluster Architecture

- Gained better insight into the cluster's architecture:
    - Head Node: Where users log in, but no GPUs are available.
    - Worker Nodes: GPUs are only available on these nodes through Slurm jobs.
  - Confirmed that Ollama needs to run on worker nodes for GPU access.
- 

## 6. Issues with Docker GPU Pass-Through

- Directly using Docker containers for Ollama with GPU pass-through was unsuccessful.
  - The cluster's GPU access is restricted, and Docker doesn't support direct GPU usage in this environment.
- 


## 7. Solution: Using Apptainer

- Identified Apptainer (formerly Singularity) as the only functional solution for GPU access on the cluster.
- Apptainer is installed on the cluster and supports GPU pass-through.
- This method allows Ollama to run with GPU acceleration via Slurm jobs.

## Prompt Generation Analysis

 Database: **crawled\_cs\_pages2**

 Total Documents: ~548 docs

 Total Time: 410 seconds | 6m 50s

Model Tested: Llama 3.2:3b

```
'grants; University of Memphis cybersecurity degree; University of '
'Memphis online courses; University of Memphis academic misconduct '
'policy; University of Memphis student conduct code; University of '
'Memphis data science program; University of Memphis computer '
'engineering; University of Memphis artificial intelligence; '
'University of Memphis IT department contact; University of '
'Memphis student organization; University of Memphis club sports; '
'University of Memphis campus map; University of Memphis parking '
'rules; University of Memphis library hours; ',
'Url': 'https://www.memphis.edu/cs/courses/syllabi/7900.pdf',
'_id': ObjectId('677ff6dcae52426f563c996'))}
-----

Processing document with _id: 677ff6dcae52426f563c998...

Generated Llama Prompts:
{'Label': 536,
 'Processing Time (s)': 0.68,
 'Prompts': 'University of Memphis computer science; University of Memphis AI '
            'research; University of Memphis admission requirements; '
            'University of Memphis data science program; Kan Yang course '
            'website; COMP 7/8998 university; University of Memphis cloud '
            'security course; MEMPHIS CS department; COMP 7998 grading policy; '
            'University of Memphis plagiarism policy; University of Memphis '
            'disability services; Cloud computing security; Internet of Things '
            'security; Attribute-based access control; Efficient search over '
            'encrypted data; Fog computing; Crowdsourcing authentication; '
            'Blockchain introduction; University of Memphis computer science '
            'courses; MEMPHIS graduate programs; University of Memphis online '
            'courses; COMP 7/8998 course description; University of Memphis '
            'research centers; Computer science department at Memphis.',
 'Url': 'https://www.memphis.edu/cs/courses/syllabi/7998.pdf',
 '_id': ObjectId('677ff6dcae52426f563c998'))}
-----

✅ Document processing job completed.
• Job completed in 410 seconds.
[jmflagg@itiger ollama]$
```



## Ollama Setup and GPU Usage on Cluster

 Check for Apptainer Installation

**which apptainer**

**/usr/bin/apptainer**

 Download and Convert Ollama Docker Image to SIF

**apptainer pull docker://ollama/ollama:latest**

- **Converts Docker image to SIF format for Apptainer compatibility.**

 Verify SIF File

**ls -lh \*.sif**

**# Example Output:**

**# -rwxr-xr-x 1 jmflagg users 1.7G Mar 26 18:44 ollama\_latest.sif**



## Cluster Management



Check Node Status

**sinfo**

- **Provides general info about available nodes.**



Detailed Node Information

**sinfo -N -o "%N %P %C %G %T %M %E"**

- Displays node-specific details like CPU, GPU availability, state, and errors.
- 

## Running Jobs on Specific Nodes

### SSH into Cluster

```
ssh jmflagg@itiger.memphis.edu
```

### Navigate to Project Directory

```
cd /project/jmflagg/ollama
```

### Submit a Job to Specific Node

```
srun --partition=bigTiger --nodelist=itiger03 --gres=gpu:1 --mem=64G  
--time=1:00:00 --pty bash
```

- Requests a GPU node for one hour with 64GB memory.

### Check Running Jobs

```
squeue -u $USER
```

### Monitor GPU Status

```
nvidia-smi
```

- Displays real-time GPU usage and memory allocation.

### Submit Job with SBATCH to specific node

```
sbatch --nodelist=itiger03 app-job.sh
```

---

## Install Ollama Natively

### Download Ollama Binary

```
curl -L https://ollama.com/download/ollama-linux-amd64.tgz -o  
ollama-linux-amd64.tgz
```

### Extract and Install

```
mkdir -p ~/ollama
```

```
tar -xzf ollama-linux-amd64.tgz -C ~/ollama
```

### Add Ollama to Path

```
export PATH=$HOME/ollama/bin:$PATH
```

### # Confirm installation

```
ollama --version
```